

# Basic Knowledge

## Axis

- 计算时，保留该轴的数量元素；0是横轴，1是纵轴

```

1 In [1]: import numpy as np
2 #生成一个3行4列的数组
3 In [2]: a = np.arange(12).reshape(3,4)
4 In [3]: a
5 Out[3]:
6 array([[ 0,  1,  2,  3],
7        [ 4,  5,  6,  7],
8        [ 8,  9, 10, 11]])
9 #axis= 0 对a的横轴进行操作，在运算的过程中其运算的方向表现为纵向运算
10 In [4]: a.sum(axis = 0)
11 Out[4]: array([12, 15, 18, 21])
12 #axis= 1 对a的纵轴进行操作，在运算的过程中其运算的方向表现为横向运算
13 In [5]: a.sum(axis = 1)
14 Out[5]: array([ 6, 22, 38])

```

## 多层感知机 (MLP: Multilayer Perceptron)

- MLP由三层或更多层非线性激活节点组成(一个输入层和一个具有一个或多个隐藏层的输出层)。由于多层互连是完全连接的，所以一层中的每个节点都以一定的权重  $w_{ij}$  连接到下一层的每个节点。
  - MLP 在感知器中进行学习，通过每次处理数据后改变连接权重，降低输出与预测结果的误差量。这是有监督学习的一个例子，通过反向传播来实现，反向传播是线性感知器中最小均方算法的推广。
- 我们可以将输出节点  $j$  的第  $n$  个数据点的误差表示为

$$e_j(n) = d_j(n) - y_j(n)$$

其中  $d$  是目标值， $y$  是由感知器预测的值。调整节点权重的方式是，尝试通过修正节点权重最小化输出的整体误差

$$\mathcal{E}(n) = \frac{1}{2} \sum_j e_j^2(n)$$

- 使用梯度下降，每个权重的修正量为

$$\Delta w_{ji}(n) = -\eta \frac{\partial \mathcal{E}(n)}{\partial v_j(n)} y_i(n)$$

其中  $y(i)$  是前一个神经元的输出， $\eta$ 是学习率。 $\eta$ 需要精心挑选，保证权重可以快速收敛而不发生震荡。

式中的导数取决于局部场  $v(j)$ 。场是变化的。很容易证明输出节点的导数可以简化为

$$-\frac{\partial \mathcal{E}(n)}{\partial v_j(n)} = e_j(n) \phi'(v_j(n))$$

其中

$$\phi'$$

是激活函数的导数,是不变的。对于隐藏节点的权重变化，分析更加困难，但是可以看出相关的导数是

$$-\frac{\partial \mathcal{E}(n)}{\partial v_j(n)} = \phi'(v_j(n)) \sum_k -\frac{\partial \mathcal{E}(n)}{\partial v_k(n)} w_{kj}(n)$$


代表输出层的第k个节点的权重变化会影响这个导数。因此，为了改变隐藏层权重，输出层权重根据激活函数的导数而改变，因此该算法代表激活函数的反向传播。

## KNN

- 寻找距离给定目标最近的节点，将该节点的结果赋予目标节点
- Distance Metric


## L1 (Manhattan) distance

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$



## L2 (Euclidean) distance

$$d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$



- 超参数：K和距离函数的选取

- 将数据分成train, validation, test三部分
- 选择在validation上表现好的超参数，在test上预测


train	validation	test
-------	------------	------

- 或者将数据打乱分组，训练集和验证集交叉排列

fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test


Useful for small datasets, but not used too frequently in deep learning

- KNN在高维空间运行会出现”维度诅咒” 的问题，那是因为在高维空间太广阔，高维空间的数据点不趋向接近另外的数据点




## Linear Classifier

- 将图片展开成一个向量



- 选择损失函数和优化方式，寻找使得loss最小的W
- 可以看作是高维的图像映射成二维的点



- 可以将线性分类器看作模板匹配；而且是将多个相同类别不同外形的特征作为模板，但对不同颜色的同类别事物鉴别较差
- 图像数据预处理：将图像像素的像素值限定到[-127, 127]，特征限定到[-1, 1]

## 损失函数

Multiclass SVM loss

and using the shorthand for the scores vector:  $s = f(x_i, W)$

the SVM loss has the form:

$$L_i = \sum_{j \neq y_i} \begin{cases} 0 & \text{if } s_{y_i} \geq s_j + 1 \\ s_j - s_{y_i} + 1 & \text{otherwise} \end{cases}$$

$$= \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

$$L = \underbrace{\frac{1}{N} \sum_i L_i}_{\text{data loss}} + \underbrace{\lambda R(W)}_{\text{regularization loss}}$$

Or expanding this out in its full form:

$$L = \frac{1}{N} \sum_i \sum_{j \neq y_i} [\max(0, f(x_i; W)_j - f(x_i; W)_{y_i} + \Delta)] + \lambda \sum_k \sum_l W_{k,l}^2$$

**Relation to Binary Support Vector Machine.** You may be coming to this class with previous experience with Binary Support Vector Machines, where the loss for the i-th example can be written as:

$$L_i = C \max(0, 1 - y_i w^T x_i) + R(W)$$

- 求导

$$j \neq y_i, w_j * x_i^T - w_{y_i} * x_i^T + \delta > 0 \Rightarrow \frac{\partial L_i}{\partial w_j} = x_i^T$$

$$j \neq y_i, w_j * x_i^T - w_{y_i} * x_i^T + \delta > 0 \Rightarrow \frac{\partial L_i}{\partial w_{y_i}} = -x_i^T$$

## Sofamax classifier (Multinomial Logistic Regression)

- Maximum Likelihood Estimation (MLE): 极大似然估计

Want to interpret raw classifier scores as **probabilities**

$$s = f(x_i; W)$$

$$P(Y = k | X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

Softmax Function

- 损失函数

Maximize probability of correct class

$$L_i = -\log P(Y = y_i | X = x_i)$$

Putting it all together:

$$L_i = -\log \left( \frac{e^{s_{y_i}}}{\sum_j e^{s_j}} \right)$$

$$L_i = -\log \left( \frac{e^{f_{y_i}}}{\sum_j e^{f_j}} \right) \quad \text{or equivalently} \quad L_i = -f_{y_i} + \log \sum_j e^{f_j}$$

- Softmax分类器最小化估计类别概率之间的交叉熵

$$H(p, q) = H(p) + D_{KL}(p||q)$$

- 相当于最小化KL散度

**Information theory view.** The cross-entropy between a "true" distribution  $p$  and an estimated distribution  $q$  is defined as:

$$H(p, q) = -\sum p(x) \log q(x)$$

**Practical issues: Numeric stability.** When you're writing code for computing the Softmax function in practice, the intermediate terms  $e^{f_{y_i}}$  and  $\sum_j e^{f_j}$  may be very large due to the exponentials. Dividing large numbers can be numerically unstable, so it is important to use a normalization trick. Notice that if we multiply the top and bottom of the fraction by a constant  $C$  and push it into the sum, we get the following (mathematically equivalent) expression:

$$\frac{e^{f_{y_i}}}{\sum_j e^{f_j}} = \frac{C e^{f_{y_i}}}{C \sum_j e^{f_j}} = \frac{e^{f_{y_i} + \log C}}{\sum_j e^{f_j + \log C}}$$

We are free to choose the value of  $C$ . This will not change any of the results, but we can use this value to improve the numerical stability of the computation. A common choice for  $C$  is to set  $\log C = -\max_j f_j$ . This simply states that we should shift the values inside the vector  $f$  so that the highest value is zero. In code:

- 求导：y是只有一个元素为1，z是输入；a是经过softmax的结果

(1) 当  $i \neq j$

$$\frac{\partial a_j}{\partial z_i} = \frac{0 - e^{z_j} e^{z_i}}{(\sum_k^n e^{z_k})^2} = -a_j a_i$$

$$\frac{\partial L}{\partial z_i} = -a_j a_i * -\frac{1}{a_j} = a_i$$

(2) 当  $i = j$

$$\frac{\partial a_j}{\partial z_j} = \frac{e^{z_j} \sum_k^n e^{z_k} - e^{z_j} e^{z_j}}{(\sum_k^n e^{z_k})^2} = a_j - a_j^2$$


$$\frac{\partial L}{\partial z_j} = (a_j - a_j^2) * -\frac{1}{a_j} = a_j - 1$$

所以,  $\frac{\partial L}{\partial \mathbf{z}} = [a_1, a_2, \dots, a_j - 1, \dots, a_n] = \mathbf{a} - \mathbf{y}$  .

## SVM vs. Softmax

- 二者区别不大；但是SVM只在乎正确的类别，其他类别损失均为0；而Softmax则会计算每个类别的损失

# Softmax vs. SVM



## Regularization and Optimization

### Regularization

- 防止模型训练过拟合，增加超参数 $\lambda$
- 减小权重 $W$ ，进行权重衰减，使网络复杂度降低

$$L(W) = \underbrace{\frac{1}{N} \sum_{i=1}^N L_i(f(x_i, W), y_i)}_{\text{Data loss: Model predictions}} + \lambda R(W)$$

**Data loss:** Model predictions should match training data

**Regularization:** Prevent the model from doing *too well* on training data

- 常见的正则化

## Simple examples

L2 regularization:  $R(W) = \sum_k \sum_l W_{k,l}^2$

L1 regularization:  $R(W) = \sum_k \sum_l |W_{k,l}|$

Elastic net (L1 + L2):  $R(W) = \sum_k \sum_l \beta W_{k,l}^2 + |W_{k,l}|$

- L2正则化的作用

$$\begin{aligned}\frac{\partial C}{\partial w} &= \frac{\partial C_0}{\partial w} + \frac{\lambda}{n} w \\ \frac{\partial C}{\partial b} &= \frac{\partial C_0}{\partial b}.\end{aligned}$$

$$\begin{aligned}w &\rightarrow w - \eta \frac{\partial C_0}{\partial w} - \frac{\eta \lambda}{n} w \\ &= \left(1 - \frac{\eta \lambda}{n}\right) w - \eta \frac{\partial C_0}{\partial w}.\end{aligned}$$

- 可以发现L2正则化项对b的更新没有影响，但是对于w的更新有影响：
- 所以  $1 - \eta \lambda / n$  小于 1，它的效果是减小 w，这也就是权重衰减（weight decay）的由来；更小的权值 w，从某种意义上说，表示网络的复杂度更低，对数据的拟合刚刚好（这个法则也叫做奥卡姆剃刀）**

## Optimization

- Numerical gradient: 数值计算，求导来计算梯度
- Analytic gradient: 梯度分析
- Always use analytic gradient, but check implementation with numerical gradient. This is called a gradient check.**
- 梯度下降需要考虑步长：后面叫做learning rate，如何选择合适的lr是一个重大的问题

### Mini-batch 梯度下降

- 小batch的梯度是目标梯度的良好近似
- 可以实现更快地收敛

### 随机梯度下降

- Stochastic Gradient Descent (SGD)
- 核心思想：迭代精细化，最开始使用随机的W，不断迭代精细化，每次迭代得到更好的W

$$\frac{df(x)}{dx} = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$


$$[f(x+h) - f(x-h)]/2h$$

$$L(W) = \frac{1}{N} \sum_{i=1}^N L_i(x_i, y_i, W) + \lambda R(W)$$


$$\nabla_W L(W) = \frac{1}{N} \sum_{i=1}^N \nabla_W L_i(x_i, y_i, W) + \lambda \nabla_W R(W)$$

- 出现问题

- 在某一方向变化剧烈，会造成收敛波动，收敛速度变慢



- 在局部最小值点（鞍点）收敛，不能达到最优值



- Our gradients come from minibatches so they can be noisy

**SGD**

# SGD

$$x_{t+1} = x_t - \alpha \nabla f(x_t)$$

```
while True:  
    dx = compute_gradient(x)  
    x -= learning_rate * dx
```

## SGD + Momentum (动量):

- 避免在鞍点收敛，尽可能在最优点收敛

# SGD+Momentum


$$v_{t+1} = \rho v_t + \nabla f(x_t)$$

$$x_{t+1} = x_t - \alpha v_{t+1}$$

```
vx = 0  
while True:  
    dx = compute_gradient(x)  
    vx = rho * vx + dx  
    x -= learning_rate * vx
```

- 两种动量


## Momentum update:



Combine gradient at current point with velocity to get step used to update weights

Nesterov, "A method of solving a convex programming problem with convergence rate  $O(1/k^2)$ ", 1983  
 Nesterov, "Introductory lectures on convex optimization: a basic course", 2004  
 Sutskever et al, "On the importance of initialization and momentum in deep learning", ICML 2013

## Nesterov Momentum



"Look ahead" to the point where updating using velocity would take us; compute gradient there and mix it with velocity to get actual update direction

## Nesterov Momentum

- 校正动量

$$v_{t+1} = \rho v_t - \alpha \nabla f(x_t + \rho v_t)$$


$$x_{t+1} = x_t + v_{t+1}$$

Change of variables  $\tilde{x}_t = x_t + \rho v_t$  and rearrange:

$$v_{t+1} = \rho v_t - \alpha \nabla f(\tilde{x}_t)$$

$$\begin{aligned} \tilde{x}_{t+1} &= \tilde{x}_t - \rho v_t + (1 + \rho)v_{t+1} \\ &= \tilde{x}_t + v_{t+1} + \rho(v_{t+1} - v_t) \end{aligned}$$

Annoying, usually we want update in terms of  $x_t, \nabla f(x_t)$



"Look ahead" to the point where updating using velocity would take us; compute gradient there and mix it with velocity to get actual update direction

## AdaGrad

- 根据每个维度的历史平方和数据增加了梯度的逐元素缩放
- G是平方和的对角矩阵

$$\theta_{t+1,i} = \theta_{t,i} - \frac{\eta}{\sqrt{G_{t,ii} + \epsilon}} \cdot g_{t,i}$$

# AdaGrad

```
grad_squared = 0
while True:
    dx = compute_gradient(x)
    grad_squared += dx * dx
    x -= learning_rate * dx / (np.sqrt(grad_squared) + 1e-7)
```

## RMSProp

- Leaky AdaGrad，解决Adagrad学习率急剧下降问题
- 使用的是指数加权平均，旨在消除梯度下降中的摆动，与Momentum的效果一样，某一维度的导数比较大，则指数加权平均就大，某一维度的导数比较小，则其指数加权平均就小，这样就保证了各维度导数都在一个量级，进而减少了摆动。允许使用一个更大的学习率 $\eta$

$$E[g^2]_t = 0.9E[g^2]_{t-1} + 0.1g_t^2.$$

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{E[g^2]_t + \epsilon}} g_t.$$

## RMSProp

```
grad_squared = 0
while True:
    dx = compute_gradient(x)
    grad_squared = decay_rate * grad_squared + (1 - decay_rate) * dx * dx
    x -= learning_rate * dx / (np.sqrt(grad_squared) + 1e-7)
```

## Adam (Almost use)

- Momentum和RMSprop结合起来的一种算法

除了像 Adadelta 和 RMSprop 一样存储了过去梯度的平方  $vt$  的指数衰减平均值，也像 momentum 一样保持了过去梯度  $mt$  的指  
数衰减平均值：

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t.$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2.$$

如果  $mt$  和  $vt$  被初始化为 0 向量，那它们就会向 0 偏置，所以做了**偏差校正**，通过计算偏差校正后的  $mt$  和  $vt$  来抵消这些偏差：

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}.$$

$$\hat{v}_t = \frac{v_t}{1 - \beta_2^t}.$$

**梯度更新规则：**

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t.$$

## Adam (full form)

```
first_moment = 0
second_moment = 0
for t in range(1, num_iterations):
    dx = compute_gradient(x)
    first_moment = beta1 * first_moment + (1 - beta1) * dx
    second_moment = beta2 * second_moment + (1 - beta2) * dx * dx
    first_unbias = first_moment / (1 - beta1 ** t)
    second_unbias = second_moment / (1 - beta2 ** t)
    x -= learning_rate * first_unbias / (np.sqrt(second_unbias) + 1e-7))
```

Momentum

Bias correction

AdaGrad / RMSProp

Bias correction for the fact that  
first and second moment  
estimates start at zero

Adam with **beta1 = 0.9**,  
**beta2 = 0.999**, and **learning\_rate = 1e-3 or 5e-4**  
is a great starting point for many models!