



CFS调度

调度器类



rq/cfs_rq/task_group/sched_entity

1. struct rq: 每个 CPU 都有一个对应的运行队列;
 2. struct cfs_rq: CFS 运行队列, 该结构中包含了 struct rb_root_cached 红黑树, 用于链接调度实体 struct sched_entity。rq 运行队列中对应了一个 CFS 运行队列, 此外, 在 task_group 结构中也会为每个 CPU 再维护一个 CFS 运行队列;
 3. struct task_struct: 任务的描述符, 包含了进程的所有信息, 该结构中的 struct sched_entity, 用于参与 CFS 的调度;
 4. struct task_group: 组调度, Linux 支持将任务分组来对 CPU 资源进行分配管理, 该结构中为系统中的每个 CPU 都分配了 struct sched_entity 调度实体和 struct cfs_rq 运行队列, 其中 struct sched_entity 用于参与 CFS 的调度;
 5. struct sched_entity: 调度实体, 这个也是 CFS 调度管理的对象了;



sched_entity

```

1 struct sched_entity {
2     struct load_weight load; /* for load-balancing */
3     struct rb_node    run_node;
4     struct list_head  group_node;
5     unsigned int      on_rq;
6
7     u64      exec_start;
8     u64      sum_exec_runtime;
9     u64      vruntime;
10    u64      prev_sum_exec_runtime;
11
12    u64      last_wakeup;
13    u64      avg_overlap;
14
15    u64      nr_migrations;
16
17    u64      start_runtime;
18    u64      avg_wakeup;
19 //...
20 #ifdef CONFIG_FAIR_GROUP_SCHED
21     struct sched_entity *parent;
22     /* rq on which this entity is (to be) queued: */
23     struct cfs_rq   *cfs_rq;
24     /* rq "owned" by this entity/group: */
25     struct cfs_rq   *my_q;
26 #endif
27 };

```

cfs_rq

```
1  /* CFS-related fields in a runqueue */
2  struct cfs_rq {
3      struct load_weight load;           //CFS运行队列的负载权重值
4      unsigned int nr_running, h_nr_running; //nr_running: 运行的调度实体数 (参与时间片计算)
5
6      u64 exec_clock;     //运行时间
7      u64 min_vruntime;   //最少的虚拟运行时间, 调度实体入队出队时需要进行增减处理
8 #ifndef CONFIG_64BIT
9     u64 min_vruntime_copy;
10 #endif
11
12     struct rb_root_cached tasks_timeline; //红黑树, 用于存放调度实体
13
14     /*
15      * 'curr' points to currently running entity on this cfs_rq.
16      * It is set to NULL otherwise (i.e when none are currently running).
17      */
18     struct sched_entity *curr, *next, *last, *skip; //分别指向当前运行的调度实体、下一个调度的调度实体、CFS运行队列中排最后的调度实体、跳过运行的调度实体
19
20 #ifdef CONFIG_SCHED_DEBUG
21     unsigned int nr_spread_over;
22 #endif
23
24 #ifdef CONFIG_SMP
25     /*
26      * CFS load tracking
27      */
28     struct sched_avg avg;             //计算负载相关
29     u64 runnable_load_sum;
30     unsigned long runnable_load_avg; //基于PELT的可运行平均负载
31 #ifdef CONFIG_FAIR_GROUP_SCHED
32     unsigned long tg_load_avg_contrib; //任务组的负载贡献
33     unsigned long propagate_avg;
34 #endif
35     atomic_long_t removed_load_avg, removed_util_avg;
36 #ifndef CONFIG_64BIT
37     u64 load_last_update_time_copy;
38 #endif
39
40 #ifdef CONFIG_FAIR_GROUP_SCHED
41     /*
42      * h_load = weight * f(tg)
```

```

43     *
44     * Where f(tg) is the recursive weight fraction assigned to this group.
45     */
46     unsigned long h_load;
47     u64 last_h_load_update;
48     struct sched_entity *h_load_next;
49 #endif /* CONFIG_FAIR_GROUP_SCHED */
50 #endif /* CONFIG_SMP */
51
52 #ifdef CONFIG_FAIR_GROUP_SCHED
53     struct rq *rq; /* cpu runqueue to which this cfs_rq is attached */ //指
      向CFS运行队列所属的CPU RQ运行队列
54
55     /*
56     * leaf cfs_rqs are those that hold tasks (lowest schedulable entity in
57     * a hierarchy). Non-leaf lrqs hold other higher schedulable entities
58     * (like users, containers etc.)
59     *
60     * leaf_cfs_rq_list ties together list of leaf cfs_rq's in a cpu. This
61     * list is used during load balance.
62     */
63     int on_list;
64     struct list_head leaf_cfs_rq_list;
65     struct task_group *tg; /* group that "owns" this runqueue */ //CFS运行队
      列所属的任务组
66
67 #ifdef CONFIG_CFS_BANDWIDTH
68     int runtime_enabled; //CFS运行队列中使用CFS带宽控制
69     u64 runtime_expires; //到期的运行时间
70     s64 runtime_remaining; //剩余的运行时间
71
72     u64 throttled_clock, throttled_clock_task; //限流时间相关
73     u64 throttled_clock_task_time;
74     int throttled, throttle_count; //throttled: 限流, throttle_count: CFS运行
      队列限流次数
75     struct list_head throttled_list; //运行队列限流链表节点, 用于添加到
      cfs_bandwidth结构中的cfttle_cfs_rq链表中
76 #endif /* CONFIG_CFS_BANDWIDTH */
77 #endif /* CONFIG_FAIR_GROUP_SCHED */
78 };

```

runtime与vruntime

1. Linux 内核默认的 sysctl_sched_latency 是 6ms，这个值用户态可设。 sched_period 用于保证可运行任务都能至少运行一次的时间间隔；

2. 当可运行任务大于 8 个的时候，`sched_period` 的计算则需要根据任务个数乘以最小调度颗粒值，这个值系统默认为 0.75ms；
3. 每个任务的运行时间计算，是用 `sched_period` 值，去乘以该任务在整个 CFS 运行队列中的权重占比；
4. 虚拟运行的时间 = 实际运行时间 * NICE_0_LOAD / 该任务的权重；


函数调用

```

1  /*
2   * delta /= w
3   */
4  static inline unsigned long
5  calc_delta_fair(unsigned long delta, struct sched_entity *se)
6  {
7      if (unlikely(se->load.weight != NICE_0_LOAD))
8          delta = calc_delta_mine(delta, NICE_0_LOAD, &se->load);
9
10     return delta;
11 }
12
13 /*
14  * The idea is to set a period in which each task runs once.
15  *
16  * When there are too many tasks (sysctl_sched_nr_latency) we have to stretch
17  * this period because otherwise the slices get too small.
18  *
19  * p = (nr <= nl) ? l : l*nr/nl
20 */
21 static u64 __sched_period(unsigned long nr_running)
22 {
23     u64 period = sysctl_sched_latency;
24     unsigned long nr_latency = sched_nr_latency;
25
26     if (unlikely(nr_running > nr_latency)) {
27         period = sysctl_sched_min_granularity;
28         period *= nr_running;
29     }
30
31     return period;
32 }
33
34 /*
35  * We calculate the wall-time slice from the period by taking a part
36  * proportional to the weight.

```

```
37     *
38     * s = p*P[w/rw]
39     */
40 static u64 sched_slice(struct cfs_rq *cfs_rq, struct sched_entity *se)
41 {
42     u64 slice = __sched_period(cfs_rq->nr_running + !se->on_rq);
43
44     for_each_sched_entity(se) {
45         struct load_weight *load;
46         struct load_weight lw;
47
48         cfs_rq = cfs_rq_of(se);
49         load = &cfs_rq->load;
50
51         if (unlikely(!se->on_rq)) {
52             lw = cfs_rq->load;
53
54             update_load_add(&lw, se->load.weight);
55             load = &lw;
56         }
57         slice = calc_delta_mine(slice, se->load.weight, load);
58     }
59     return slice;
60 }
61
62 /*
63 * We calculate the vruntime slice of a to be inserted task
64 *
65 * vs = s/w
66 */
67 static u64 sched_vslice(struct cfs_rq *cfs_rq, struct sched_entity *se)
68 {
69     return calc_delta_fair(sched_slice(cfs_rq, se), se);
70 }
```



以 Task A 为例: $\text{sched_slice} = 6\text{ms} * (1586 * 801000) \gg 32$

$(\text{delta_exec} * (\text{weight} * \text{lw->inv_weight})) \gg \text{WMULT_SHIFT}$



| Task A nice = -2 weight=1586 inv_weight=2708050 | Task B nice = -1 weight=1277 inv_weight=3363326 | Task C nice = 0 weight=1024 inv_weight=4194304 | Task D nice = 1 weight=820 inv_weight=523376 | Task E nice = 2 weight=655 inv_weight=6557202 |
|--|--|---|---|--|
|--|--|---|---|--|

累加各个Task的weight得到 `lw.weight=5362`
`lw.inv_weight=0xFFFFFFFF / lw.weight = 801000`

`sysctl_sched_latency = 6ms`

CFB调度 tick

- task_tick_fair->entity->tick->update_curr->_update_curr



```

1  /*
2   * scheduler tick hitting a task of our scheduling class:
3   */
4  static void task_tick_fair(struct rq *rq, struct task_struct *curr, int
queued)

```

```
5  {
6      struct cfs_rq *cfs_rq;
7      struct sched_entity *se = &curr->se;
8
9      for_each_sched_entity(se) {
10         cfs_rq = cfs_rq_of(se);
11         entity_tick(cfs_rq, se, queued);
12     }
13 }
14
15 static void
16 entity_tick(struct cfs_rq *cfs_rq, struct sched_entity *curr, int queued)
17 {
18     /*
19      * Update run-time statistics of the 'current'.
20      */
21     update_curr(cfs_rq);
22
23 #ifdef CONFIG_SCHED_HRTICK
24     /*
25      * queued ticks are scheduled to match the slice, so don't bother
26      * validating it and just reschedule.
27      */
28     if (queued) {
29         resched_task(rq_of(cfs_rq)->curr);
30         return;
31     }
32     /*
33      * don't let the period tick interfere with the hrtick preemption
34      */
35     if (!sched_feat(DOUBLE_TICK) &&
36         hrtimer_active(&rq_of(cfs_rq)->hrtick_timer))
37         return;
38 #endif
39
40     if (cfs_rq->nr_running > 1 || !sched_feat(WAKEUP_PREEMPT))
41         check_preempt_tick(cfs_rq, curr);
42 }
43
44 /*
45  * Update the current task's runtime statistics. Skip current tasks that
46  * are not in our scheduling class.
47 */
48 static inline void
49 __update_curr(struct cfs_rq *cfs_rq, struct sched_entity *curr,
50             unsigned long delta_exec)
51 {
```

```

52     unsigned long delta_exec_weighted;
53
54     schedstat_set(curr->exec_max, max((u64)delta_exec, curr->exec_max));
55
56     curr->sum_exec_runtime += delta_exec;
57     schedstat_add(cfs_rq, exec_clock, delta_exec);
58     delta_exec_weighted = calc_delta_fair(delta_exec, curr);
59
60     curr->vruntime += delta_exec_weighted;
61     update_min_vruntime(cfs_rq);
62 }
63
64
65 static void update_curr(struct cfs_rq *cfs_rq)
66 {
67     struct sched_entity *curr = cfs_rq->curr;
68     u64 now = rq_of(cfs_rq)->clock;
69     unsigned long delta_exec;
70
71     if (unlikely(!curr))
72         return;
73
74     /*
75      * Get the amount of time the current task was running
76      * since the last time we changed load (this cannot
77      * overflow on 32 bits):
78      */
79     delta_exec = (unsigned long)(now - curr->exec_start);
80     if (!delta_exec)
81         return;
82
83     __update_curr(cfs_rq, curr, delta_exec);
84     curr->exec_start = now;
85
86     if (entity_is_task(curr)) {
87         struct task_struct *curtask = task_of(curr);
88
89         trace_sched_stat_runtime(curtask, delta_exec, curr->vruntime);
90         cpuacct_charge(curtask, delta_exec);
91         account_group_exec_runtime(curtask, delta_exec);
92     }
93 }

```

任务出队入队

- 当任务进入可运行状态时，需要将调度实体放入到红黑树中，完成入队操作；

2. 当任务退出可运行状态时，需要将调度实体从红黑树中移除，完成出队操作；

3. CFS 调度器，使用 `enqueue_entity` 函数将任务入队到 CFS 队列，使用 `dequeue_entity` 函数将任务从 CFS 队列中出队操作

```
1 static void
2 enqueue_entity(struct cfs_rq *cfs_rq, struct sched_entity *se, int flags)
3 {
4     /*
5      * Update the normalized vruntime before updating min_vruntime
6      * through calling update_curr().
7      */
8     if (!(flags & ENQUEUE_WAKEUP) || (flags & ENQUEUE_MIGRATE))
9         se->vruntime += cfs_rq->min_vruntime;
10
11    /*
12     * Update run-time statistics of the 'current'.
13     */
14    update_curr(cfs_rq);
15    account_entity_enqueue(cfs_rq, se);
16
17    if (flags & ENQUEUE_WAKEUP) {
18        place_entity(cfs_rq, se, 0);
19        enqueue_sleeper(cfs_rq, se);
20    }
21
22    update_stats_enqueue(cfs_rq, se);
23    check_spread(cfs_rq, se);
24    if (se != cfs_rq->curr)
25        __enqueue_entity(cfs_rq, se);
26 }
27
28 /*
29 * Enqueue an entity into the rb-tree:
30 */
31 static void __enqueue_entity(struct cfs_rq *cfs_rq, struct sched_entity *se)
32 {
33     struct rb_node **link = &cfs_rq->tasks_timeline.rb_node;
34     struct rb_node *parent = NULL;
35     struct sched_entity *entry;
36     s64 key = entity_key(cfs_rq, se);
37     int leftmost = 1;
38
39     /*
40      * Find the right place in the rbtree:
41      */
```

```

42     while (*link) {
43         parent = *link;
44         entry = rb_entry(parent, struct sched_entity, run_node);
45         /*
46          * We dont care about collisions. Nodes with
47          * the same key stay together.
48         */
49         if (key < entity_key(cfs_rq, entry)) {
50             link = &parent->rb_left;
51         } else {
52             link = &parent->rb_right;
53             leftmost = 0;
54         }
55     }
56
57     /*
58      * Maintain a cache of leftmost tree entries (it is frequently
59      * used):
60     */
61     if (leftmost)
62         cfs_rq->rb_leftmost = &se->run_node;
63
64     rb_link_node(&se->run_node, parent, link);
65     rb_insert_color(&se->run_node, &cfs_rq->tasks_timeline);
66 }

```

寻找下一个任务

- `__pick_next_entity ()`：返回红黑树最左侧节点，该节点是vruntime最小的节点
- 使用`rb_leftmost`缓存最小vruntime节点

```

1  static struct sched_entity *__pick_next_entity(struct cfs_rq *cfs_rq)
2  {
3     struct rb_node *left = cfs_rq->rb_leftmost;
4
5     if (!left)
6         return NULL;
7
8     return rb_entry(left, struct sched_entity, run_node);
9 }

```